



Today & Next week

(SQL, SQL, SQL)

Phase I: <u>Intuition</u> for SQL (1st half of today)

Basic Relational model (aka tables)

Example SQL (exploring real datasets)

Phase II: Formal description

SQL concepts we'll study (similar to Python map-reduce)

Schemas, Query structure of SELECT-FROM-WHERE, JOINs, etc



Data independence

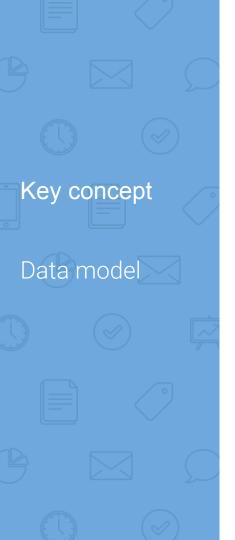
1. Can we add a new column or attribute without rewriting the application?

Logical Data Independence

Protection from changes in the Logical Structure of the data

2. Do you need to care which disks/machines are the data stored on?

<u>Physical Data Independence</u> Protection from Physical Layout Changes



Relational model (aka tables)
Simple and most popular
Elegant algebra (E.F. Codd et al)

<u>Data model</u>:

Organizing principle of data + operations

Every Relation has a Schema

Logical Schema: describes types, names

Physical Schema: describes data layout

Virtual Schema (Views): derived tables

Schema:

Describes blueprint of table (s)



Set algebra (reminder)

List: [1, 1, 2, 3]

Set: {1, 2, 3}

Multiset: {1, 1, 2, 3}

A <u>multiset</u> is an unordered list (or: a set with multiple duplicate instances allowed)

UNIONs

Set: $\{1, 2, 3\} \cup \{2\} = \{1, 2, 3\}$

<u>Multiset</u>: {1, 1, 2, 3} U { 2 } = { 1, 1, 2, 2, 3 }

Cross-product

```
{1, 1, 2, 3} * { y, z } =

{ <1, y>, <1, y>, <2, y>, <3, y>

<1, z>, <1, z>, <2, z>, <3, z>

}
```

Python operating on Lists [reminder]

BASIC TYPES

Int, long int string ...

MAP + FILTER

map(function, list of inputs)

filter(function, list of inputs)

- Map applies function to input list
- Filter returns sub-list that satisfies a filter condition

REDUCE/AGGREGATE

reduce (...)

- Reduce runs a computation on a list and returns a result. E.g., SUM, MAX, MIN

For review, check out your favorite python tutorial (e.g, https://book.pythontips.com/en/latest/map_filter.html

SQL Queries on Tables (Lists of rows)

BASIC TYPES

Int32, int64 Char[n] ... Float32, float64

MAP + FILTER

Single table query

SELECT c1, c2 FROM T WHERE condition;

Multi table JOIN

SELECT c1, c2 FROM T1, T2 WHERE condition;

REDUCE/AGGREGATE

SELECT SUM(c1*c2) FROM T WHERE condition GROUP BY c3;

Map-Filter-Reduce pattern: Same simple/powerful idea in MapReduce, Hadoop, Spark, etc.



1. SQL introduction & schema definitions

- 2. Basic single-table queries
- 3. Multi-table queries



SQL Introduction

- SQL is a standard language for querying and manipulating data
- SQL is a very high-level programming language
 This works because it is optimized well!
- Many standards out there:
 ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3),

NB: Probably the world's most successful **parallel** programming language (multicore?)

<u>SQL</u> stands for<u>S</u>tructured<u>Q</u>uery<u>L</u>anguage



Data Types in SQL

Atomic types for columns:

Characters: CHAR(20), VARCHAR(50)

Numbers: INT, BIGINT, SMALLINT, FLOAT

Others: MONEY, DATETIME...

(Most SQL dialects support **record** types, e.g. json-like. Not important for cs145.)



Tables

Product

A <u>relation</u> or <u>table</u> is a multiset of rows, with columns.

The **schema** of a table is the table name, its columns, and their types.

PName	Price	Manuf
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

Each <u>row</u> (or <u>tuple</u> or <u>record)</u> has attributes

The number of tuples is the **cardinality** of the relation

Each <u>column</u> (or <u>attribute</u>) has a **type**The number of columns is the <u>arity</u> of the relation.



Key constraints

A **key** is a **minimal subset of columns** that acts as a unique identifier for tuples in a relation

i.e. if two tuples agree on the values of the key, then they must be the same tuple!

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)

Design choices?

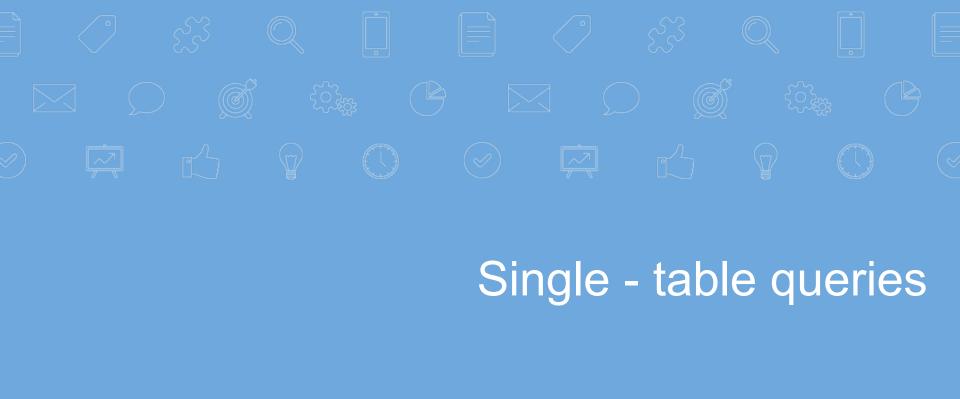
- 1. Which would you select as a key?
- 2. Is a key always guaranteed to exist?
- 3. Can we have more than one key?



Declaring Schema

```
CREATE TABLE Product (
Pname CHAR(20),
Manufacturer VARCHAR(50),
price float,
Category VARCHAR(50),
PRIMARY KEY (Pname, Manufacturer))
```

Product(Pname: string, Price: float, Category: string, Manufacturer: string)





- 1. The SFW query
- 2. Other useful operators: LIKE, DISTINCT, ORDER BY



SQL Query

Basic form (there are many many more bells and whistles)

SELECT <attributes>

FROM <one or more relations>

WHERE <conditions>

Call this a **SFW** query.



Simple SQL Query: Selection

Selection is the operation of filtering a relation's tuples on some condition

PName	Price	Category	Manuf
Gizmo	\$19.99	Gadgets	GWorks
Powergizmo	\$29.99	Gadgets	GWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

SELECT*

FROM Product

WHERE Category = 'Gadgets'



PName	Price	Category	Manuf
Gizmo	\$19.99	Gadgets	GWorks
Powergizmo	\$29.99	Gadgets	GWorks



Simple SQL Query: Projection

Projection is the operation of producing an output table with tuples that have a subset of their prior attributes

PName	Price	Category	Manuf
Gizmo	\$19.99	Gadgets	GWorks
Powergizmo	\$29.99	Gadgets	GWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

SELECT Pname, Price, Manufacturer

FROM Product

WHERE Category = 'Gadgets'



PName	Price	Manuf
Gizmo	\$19.99	GWorks
Powergizmo	\$29.99	GWorks



Notation

Input Schema

Product(<u>PName</u>, Price, Category, <u>Manufacturer</u>)

SELECT Pname, Price, Manufacturer FROM Product
WHERE Category = 'Gadgets'

Output Schema

Answer(PName, Price, Manufacturer)



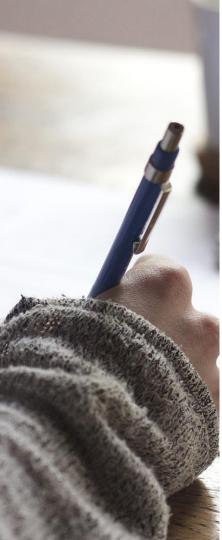
LIKE: Simple String Pattern Matching

SELECT*

FROM Products

WHERE PName LIKE '%gizmo%'

- s LIKE p: pattern matching on strings
- p may contain two special symbols:
 - % = any sequence of characters
 - o _ = any single character



DISTINCT: Eliminating Duplicates

SELECT DISTINCT Category
FROM Product



Category

Gadgets

Photography

Household

Versus

SELECT Category
FROM Product



Category

Gadgets

Gadgets

Photography

Household



ORDER BY: Sorting the Results

SELECT PName, Price, Manufacturer

FROM Product

WHERE Category='gizmo' AND Price > 50

ORDER BY Price, PName

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.



A Few Details

- 1. SQL **commands** are case insensitive:
 - a. SELECT = Select, Product = product
- 2. **Values** are **not**: 'Seattle' vs 'seattle'

- 3. Use single quotes for constants:
 - a. 'abc' best practice (versus "abc" with mixed support)
- 4. To say "don't know the value" we use NULL
 - a. Common, but annoying (more detail later)
 - b. E.g., Student GPA in 1st quarter = NULL, not zero





- 1. Foreign key constraints
- 2. Joins: basics
- 3. Joins: SQL semantics



Foreign Key constraints

Suppose we have the following schema :

Students(<u>sid</u>: *string*, name: *string*, gpa: *float*)

Enrolled(<u>student_id</u>: *string*, <u>cid</u>: *string*, grade: *string*)

And we want to impose the following constraint:
 a student must exist in the Students table to enroll in a class

Students

sid	name	gpa
102	Bob	3.9
123	Mary	3.8

Enrolled

Student _id	cid	grade
123	564	А
123	537	A+

We say that student_id is a foreign key that refers to Students



Declaring Foreign Keys

```
Students(<u>sid</u>: string, name: string, gpa: float)
Enrolled(student id: string, cid: string, grade: string)
CREATE TABLE Enrolled (
 student id CHAR(20),
 cid CHAR(20),
 grade CHAR(10),
 PRIMARY KEY (student id, cid),
 FOREIGN KEY (student id) REFERENCES Students(sid)
```



Foreign Keys and update operations

Students(<u>sid</u>: *string*, name: *string*, gpa: *float*)

Enrolled(<u>student_id</u>: *string*, <u>cid</u>: *string*, grade: *string*)

- What if we insert a tuple into Enrolled, but no corresponding student? INSERT is rejected (foreign keys are <u>constraints</u>)!
- What if we delete a student? Design choices
 - 1. Disallow the delete
 - 2. Remove all of the courses for that student
 - 3. SQL allows a third via NULL (not yet covered)

DBA chooses



Keys and Foreign Keys

Company

<u>CName</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

What is a foreign key vs. a key here?

Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi



1. SQL introduction & schema definitions

- 2. Basic single-table queries
- 3. Multi-table queries

Preview

SQL queries

sqltutorial.org/sql-cheat-sheet

SQL CHEAT SHEET http://www.sqltutorial.org

OUERYING DATA FROM A TABLE

SELECT cl, c2 FROM t;

Query data in columns c1, c2 from a table

SELECT * FROM t:

Query all rows and columns from a table

SELECT c1, c2 FROM t

WHERE condition;

Query data and filter rows with a condition

SELECT DISTINCT cl FROM t

WHERE condition;

Query distinct rows from a table

SELECT c1, c2 FROM t

ORDER BY cl ASC [DESC];

Sort the result set in ascending or descending order

SELECT c1, c2 FROM t

ORDER BY cl

LIMIT n OFFSET offset;

Skip offset of rows and return the next n rows

SELECT c1, aggregate(c2)

FROM t

GROUP BY cl:

Group rows using an aggregate function

SELECT cl, aggregate(c2)

FROM t

GROUP BY cl

HAVING condition:

Filter groups using HAVING clause

OUERYING FROM MULTIPLE TABLES

SELECT c1, c2

FROM t1

INNER JOIN t2 ON condition;

Inner join t1 and t2

SELECT c1, c2

FROM t1

LEFT JOIN t2 ON condition;

Left join t1 and t1

SELECT c1, c2

FROM t1

RIGHT JOIN t2 ON condition:

Right join t1 and t2

SELECT c1, c2

FROM t1

FULL OUTER JOIN t2 ON condition;

Perform full outer join

SELECT c1, c2

FROM t1

CROSS JOIN t2;

Produce a Cartesian product of rows in tables

SELECT cl. c2

FROM t1, t2;

Another way to perform cross join

SELECT cl. c2

FROM t1 A

INNER JOIN t2 B ON condition;

Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT cl. c2 FROM tl

UNION [ALL]

SELECT cl, c2 FROM t2:

Combine rows from two queries

SELECT c1, c2 FROM t1

INTERSECT

SELECT c1, c2 FROM t2;

Return the intersection of two queries

SELECT c1, c2 FROM t1

MINUS

SELECT c1, c2 FROM t2;

Subtract a result set from another result set

SELECT c1, c2 FROM t1

WHERE cl [NOT] LIKE pattern;

Query rows using pattern matching %, _

SELECT cl, c2 FROM t

WHERE cl [NOT] IN value list;

Query rows in a list

SELECT cl, c2 FROM t

WHERE cl BETWEEN low AND high;

Query rows between two values

SELECT cl. c2 FROM t

WHERE CLIS [NOT] NULL:

Check if values in a table is NULL or not