



# Transactions

# Exam Grades

## Stats (Out of 90+2)

- Median: 76.5, Mean: 73.9, High: 92, Standard Deviation: 10.5
- Regrade requests open for **1 week** (next Tuesday)

## Feedback Themes

- TAs were ultra-responsive during the “24 hours”
- Test was too long. We’ll recalibrate for Finals
- Most liked 24 hour flex time



# SQL Writes

### MANAGING TABLES

```
CREATE TABLE t (  
  id INT PRIMARY KEY,  
  name VARCHAR NOT NULL,  
  price INT DEFAULT 0  
);
```

Create a new table with three columns

```
DROP TABLE t;
```

Delete the table from the database

```
ALTER TABLE t ADD column;
```

Add a new column to the table

```
ALTER TABLE t DROP COLUMN c;
```

Drop column c from the table

```
ALTER TABLE t ADD constraint;
```

Add a constraint

```
ALTER TABLE t DROP constraint;
```

Drop a constraint

```
ALTER TABLE t1 RENAME TO t2;
```

Rename a table from t1 to t2

```
ALTER TABLE t1 RENAME c1 TO c2;
```

Rename column c1 to c2

```
TRUNCATE TABLE t;
```

Remove all data in a table

### USING SQL CONSTRAINTS

```
CREATE TABLE t(  
  c1 INT, c2 INT, c3 VARCHAR,  
  PRIMARY KEY (c1,c2)  
);
```

Set c1 and c2 as a primary key

```
CREATE TABLE t1(  
  c1 INT PRIMARY KEY,  
  c2 INT,  
  FOREIGN KEY (c2) REFERENCES t2(c2)  
);
```

Set c2 column as a foreign key

```
CREATE TABLE t(  
  c1 INT, c1 INT,  
  UNIQUE(c2,c3)  
);
```

Make the values in c1 and c2 unique

```
CREATE TABLE t(  
  c1 INT, c2 INT,  
  CHECK(c1 > 0 AND c1 >= c2)  
);
```

Ensure c1 > 0 and values in c1 >= c2

```
CREATE TABLE t(  
  c1 INT PRIMARY KEY,  
  c2 VARCHAR NOT NULL  
);
```

Set values in c2 column not NULL

### MODIFYING DATA

```
INSERT INTO t(column_list)  
VALUES(value_list);
```

Insert one row into a table

```
INSERT INTO t(column_list)  
VALUES (value_list),  
      (value_list), ....;
```

Insert multiple rows into a table

```
INSERT INTO t1(column_list)  
SELECT column_list  
FROM t2;
```

Insert rows from t2 into t1

```
UPDATE t  
SET c1 = new_value;
```

Update new value in the column c1 for all rows

```
UPDATE t  
SET c1 = new_value,  
    c2 = new_value  
WHERE condition;
```

Update values in the column c1, c2 that match the condition

```
DELETE FROM t;
```

Delete all data in a table

```
DELETE FROM t  
WHERE condition;
```

Delete subset of rows in a table

# SQL Writes

```
UPDATE Product  
SET Price = Price - 1.99  
WHERE pname = 'Gizmo'
```

```
INSERT INTO SmallProduct(name, price)  
SELECT pname, price  
FROM Product  
WHERE price <= 0.99
```

```
DELETE Product  
WHERE price <= 0.99
```




# How?

## Example Game App

DB v0

(Recap lectures)



- Q1: 1000 users/sec?
- Q2: Offline?
- Q3: Support v1, v1' versions?

App designer

- Q7: How to model/evolve game data?
- Q8: How to scale to millions of users?
- Q9: When machines die, restore game state gracefully?

Systems designer

- Q4: Which user cohorts?
- Q5: Next features to build?
- Experiments to run?
- Q6: Predict ads demand?

Product/Biz designer

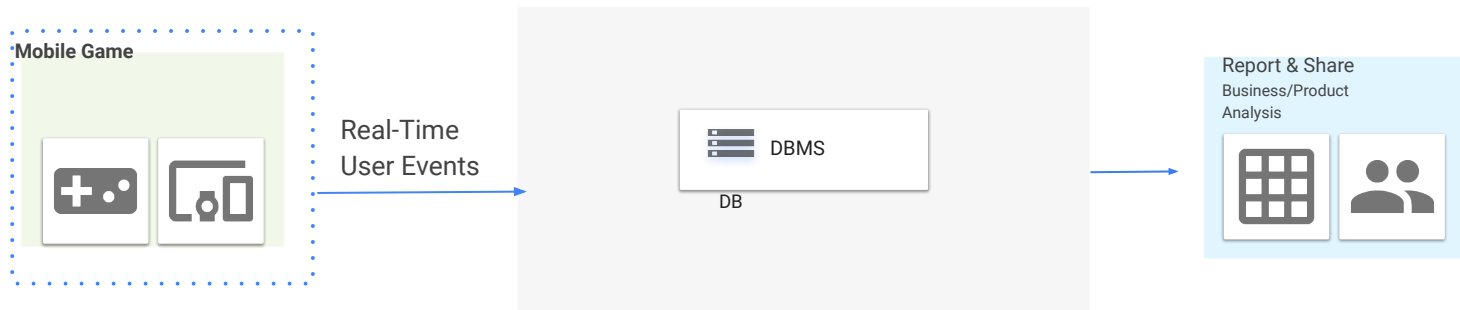


# How?

## Example Game App

DB v0

(Recap lectures)



Q1: 1000 users/sec?

Q2: Offline?

Q3: Support v1, v1' versions?

Q7: How to model/evolve game data?

Q8: How to scale to millions of users?

Q9: When machines crash, restore game state gracefully?

Q4: Which user cohorts?

Q5: Next features to build?

Experiments to run?

Q6: Predict ads demand?

App designer

Systems designer

Product/Biz designer



## Today's Lecture


1. Why Transactions?
2. Transactions
3. Properties of Transactions: ACID
4. Logging



Example

Unpack  
ATM DB:

Transaction



Read Balance  
Give money  
Update Balance

vs

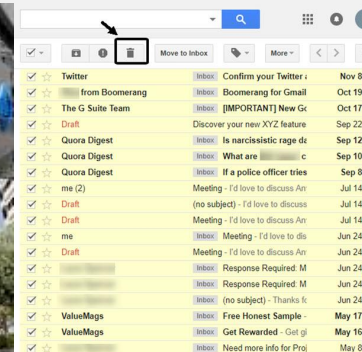
Read Balance  
Update Balance  
Give money



Visa does > 60,000 TXNs/sec with users & merchants

Want your 4\$ Starbucks transaction to wait for a stranger's 10k\$ bet in Las Vegas ?  
⇒ Transactions can (1) be quick or take a long time, (2) unrelated to you





Transactions are at the core of

- payment, stock market, banks, ticketing
- Gmail, Google Docs (e.g., multiple people editing)

# Example

## Monthly bank interest transaction

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 100M bank accounts

Takes 24 hours to run

**UPDATE** Money  
**SET** Balance = Balance \* 1.1

# Example

Monthly  
bank  
interest  
transaction

## Performance

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

### Cost to update all data

100M bank accounts → 100M seeks? (worst case)

(@10 msec/seek, that's 1 million secs)



Problem1: SLOW :(

# Example

Monthly  
bank  
interest  
transaction

With crash

Money

Account	...	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@10:45 am)

Account	...	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		...
30108		-110
40008		110
50002		22

??

?? Did T-Monthly-423 complete?  
?? Which tuples are bad?

??

Case1: T-Monthly-423 crashed  
Case2: T-Monthly-423 completed  
4002 deposited 20\$ at 10:45 am

## 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 100M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am

Problem 2: Wrong :(

# 15


Big Scale

Roadmap

Primary data structures/algorithms

LOGS

LOCKS



?????



## Today's Lecture

1. Why Transactions?
2. Properties of Transactions: ACID
3. Logging



# Transactions: Basic Definition

A transaction ("TXN") is a sequence of one or more *operations* (reads or writes) which reflects *a single real-world transition*.

In the real world, a TXN either happened completely or not at all (e.g., you withdrew 100\$ from bank. Or not.)

```
START TRANSACTION
UPDATE Product
SET Price = Price - 1.99
WHERE pname = 'Gizmo'
COMMIT
```

# Transactions in SQL

- In “ad-hoc” SQL, each statement = one transaction
- In a program, multiple statements can be grouped together as a transaction

```
START TRANSACTION
```

```
UPDATE Bank SET amount = amount - 100
```

```
WHERE name = 'Bob'
```

```
UPDATE Bank SET amount = amount + 100
```

```
WHERE name = 'Joe'
```

```
COMMIT
```

# Motivation for Transactions

Group user actions (reads & writes) into *Transactions* helps with two goals:

1. Recovery & Durability: Keep the data consistent and durable.  
*Despite system crashes, user canceling TXN part way, etc.*

This lecture!

**Idea:** Use **LOGS**. Support to “commit” or “rollback” TXNs

2. Concurrency: Get better performance by parallelizing TXNs  
*without creating ‘bad data.’ Despite slow disk writes and reads.*

Next lecture

**Idea:** Use **LOCKS**. Run several user TXNs concurrently.

# Example 1: Protection against crashes / aborts

Scenario: Make a CheapProducts table, from a Products table

Client 1:

```
INSERT INTO CheapProduct(name, price)
SELECT pname, price
FROM Product
WHERE price <= 0.99
```

**Crash / abort!**

```
DELETE Product
WHERE price <=0.99
```

What goes wrong?

Client 1:

START TRANSACTION

INSERT INTO CheapProduct(name, price)

SELECT pname, price

FROM Product

WHERE price <= 0.99

DELETE Product

WHERE price <=0.99

COMMIT

Now we'd be fine! We'll see how / why this lecture

## Example 2: Multiple users: single statements

Client 1: [at 10:01 am]

```
UPDATE Product  
SET Price = Price - 1.99  
WHERE pname = 'Gizmo'
```

Client 2: [at 10:01 am]

```
UPDATE Product  
SET Price = Price*0.5  
WHERE pname='Gizmo'
```

Two managers attempt to discount products ***at same time*** -

What could go wrong?

Client 1: START TRANSACTION

UPDATE Product

SET Price = Price – 1.99

WHERE pname = 'Gizmo'

COMMIT

Client 2: START TRANSACTION

UPDATE Product

SET Price = Price\*0.5

WHERE pname='Gizmo'

COMMIT

Now works like a charm- we'll see how / why next lecture...



### 3. Properties of Transactions



What you will  
learn about in  
this section

1. Atomicity
2. Consistency
3. Isolation
4. Durability

A close-up photograph of a hand holding a blue pen, poised to write on a piece of paper. The hand is wearing a grey, textured sweater. The background is blurred, showing a desk and a laptop.

# ACID: Atomicity

- TXN is all or nothing
  - *Commits*: all the changes are made
  - *Aborts*: no changes are made



# ACID: Consistency

- The tables must always satisfy user-specified *integrity constraints*
  - E.g., Account number is unique, Sum of *debits* and of *credits* is 0
- How consistency is achieved:
  - Programmer writes a TXN to go from one consistent state to a consistent state
  - *System* makes sure that the TXN is atomic (e.g., if EXCEPTION, rolls back)



# ACID: Isolation

- A TXN executes **concurrently** with other TXNs
- Effect of TXNs is the same as TXNs running one after another

Conceptually,

- similar to OS “sandboxes”
- E.g. TXNs can’t observe each other’s “partial updates”



# ACID: Durability

- The effect of a TXN must **persist** after the TXN
  - And after the whole program has terminated
  - And even if there are power failures, crashes, etc.
- ⇒ Write data to durable IO (e.g., disk)



# ACID Summary

- **A**tomic
  - State shows either all the effects of TXN, or none of them
- **C**onsistent
  - TXN moves from a state where integrity holds, to another where integrity holds
- **I**solated
  - Effect of TXNs is the same as TXNs running one after another
- **D**urable
  - Once a TXN has committed, its effects remain in the database

# A Note: ACID is one popular option!

- Many debates over ACID, both **historically** and **currently**
- Some “NoSQL” DBMSs relax ACID
- In turn, now “NewSQL” reintroduces ACID compliance to NoSQL-style DBMSs...

⇒ Usually, depends on what consistency and performance your application needs



ACID is an extremely important & successful paradigm,  
but still debated!



## 4. Atomicity & Durability via Logging



# Conceptual Idea: Trip to Europe

Drag a column header here to group by that column

Done	Name	Category
<input type="checkbox"/>	Download Travel Checklist	Preparation
<input type="checkbox"/>	Take an umbrella for traveling to	Packing
<input type="checkbox"/>	Take a French dictionary for traveling	Packing
<input type="checkbox"/>	Take sunglasses for traveling to	Packing
<input type="checkbox"/>	Walk up the Eiffel Tower	Paris
<input type="checkbox"/>	Have a coffee at a cafe	Paris
<input type="checkbox"/>	Go and see the Queen	London
<input type="checkbox"/>	Listen to Big Ben rings	London
<input type="checkbox"/>	Drink tea at 5 pm	London
<input type="checkbox"/>	Visit Opera House	Odessa
<input type="checkbox"/>	Go down the Petyomin Stairs	Odessa

1. Make TODO list. Buy tickets



2. Actual Visit

(Much longer than buying tickets)

Big Idea

# LOGS!

(aka TODO/  
ledger)

Recall (on disks)

- ▷ Sequential reads FASTER than random reads
- ▷ Sequential writes (aka “appends”) FASTER than random writes

## Big Idea: LOGs (or log files or ledger)

- ▷ Any value that changes? Append to LOG!
  - LOG is a compact “todo” list of data updates
- ▷ Intuition:
  - Data pages: (a) Update in RAM (fast) (b) Update on disk later (slow)
  - LOGs: ( c) Append “todo” in LOGs and (d) control when you Flush LOGs to disk

Many kinds of LOGs. We'll study a few key ones!

What you will  
learn about in  
this section

1. How to make/use LOGs?
2. How to make it fast? (Mess with memory and disk)



# Basic Idea: (Physical) Logging

Idea:

- Log consists of an ordered list of Update Records
- Log record contains UNDO information for every update!  
<TransactionID, &reference, old value, new value>  
(e.g., key)

What DB does?

- Owns the log “service” for all applications/transactions.
- Appends to log. **Flush** when necessary — force writes to disk

This is sufficient to UNDO any transaction!

# Example

Monthly  
bank  
interest  
transaction

## Full run

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

WA Log (@4:29 am day+1)

T-Monthly-423	<b>START TRANSACTION</b>		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352
T-Monthly-423	...	...	...
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	<b>COMMIT</b>		

Update  
Records

Commit  
Record

### 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 100M bank accounts

Takes 24 hours to run

START TRANSACTION

UPDATE Money

SET Amt = Amt \* 1.10

COMMIT

# Example

Monthly  
bank  
interest  
transaction

With crash

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@10:45 am)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		...
30108		-110
40008		110
50002		22

WA Log (@10:29 am)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	...	...	...
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352

??

??

??

??

## TXN 'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 100M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am

Did T-Monthly-423 complete?

Which tuples are bad?

Case1: T-Monthly-423 was crashed

Case2: T-Monthly-423 completed. 4002 deposited 20\$ at 10:45 am

Can you infer from RED log records?

# Example

## Monthly bank interest transaction

## Recovery

Money (@10:45 am)

Account	...	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		
30108		-110
40008		110
50002		22

Money (after recovery)

Account	...	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

WA Log (@10:29 am)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	...	...	...
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22

System recovery (after 10:45 am)


1. Rollback uncommitted transactions
  - Restore old values from WAL Log (if any)
  - Notify developers about aborted TXN
2. Redo Recent transactions (w/ new values)
3. Back in business; Redo (any pending) transactions

What you will  
learn about in  
this section

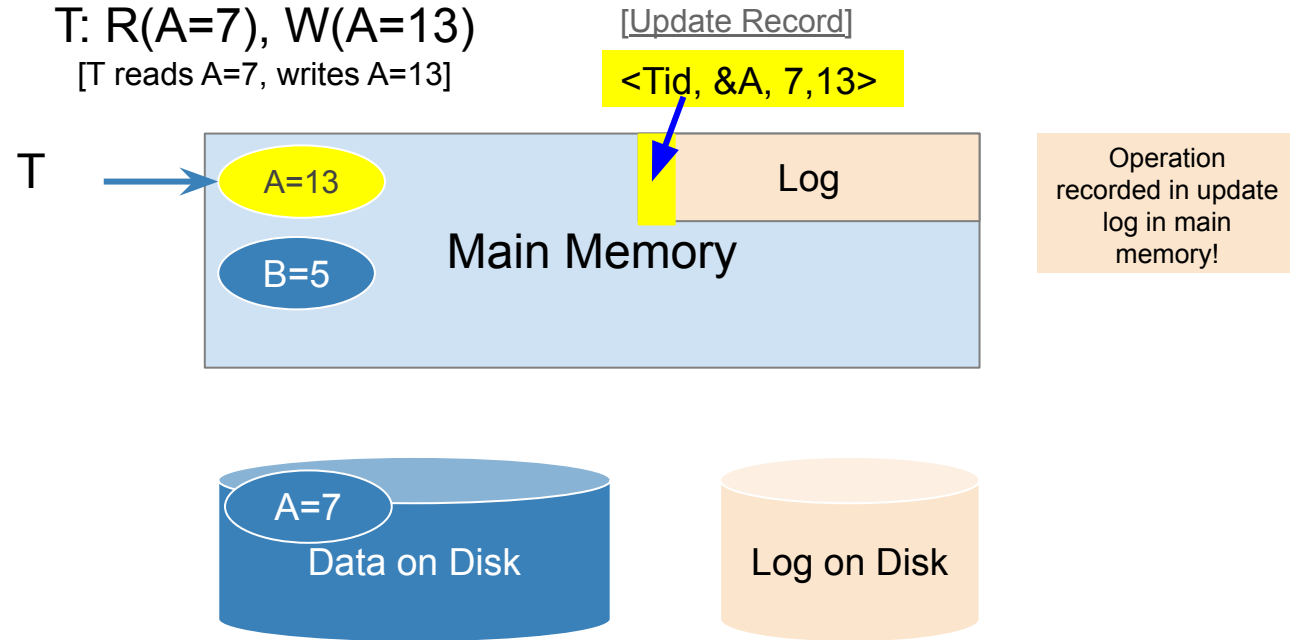
1. How to make/use LOGs?
2.  $\Rightarrow$  How to make it fast? (Mess with memory and disk)



# A picture of logging



# A picture of logging





# Why do we need logging for atomicity?

- Could we just write TXN updates to disk **only** once whole TXN complete?
  - Then, if abort / crash and TXN not complete, it has no effect- atomicity!
  - *With unlimited memory and time, this could work...*
- ⇒ We **need to log partial results of TXNs** because of:
  - Memory constraints (e.g. , billions of updates)
  - Time constraints (what if one TXN takes very long?)

We need to write partial results to disk!  
...And so we need a **LOG** to (maybe) **undo** these partial results!

# What is the correct way to LOG to disk?

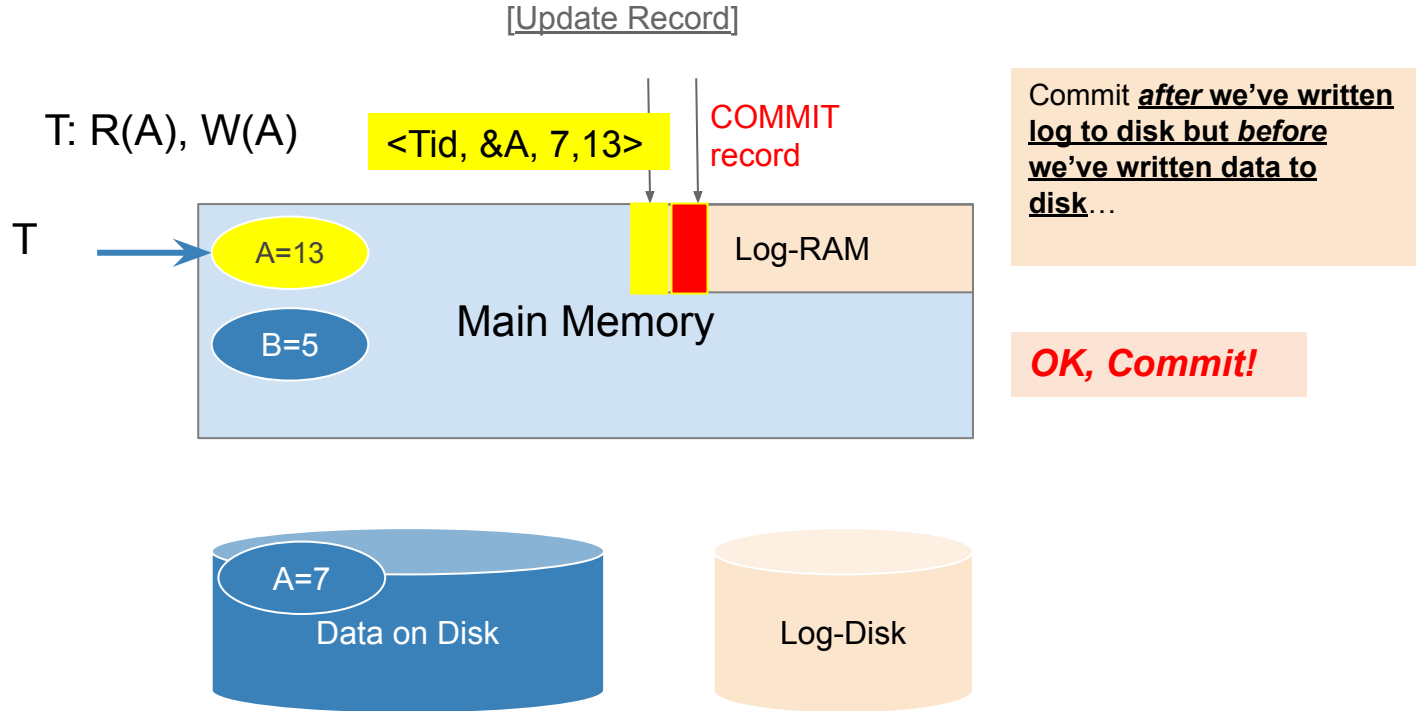
- We'll look at the *Write-Ahead Logging (WAL)* protocol
- We'll see why it works by looking at other protocols which are incorrect!

Remember: Key idea is to ensure durability  
*while* maintaining our ability to “undo”!

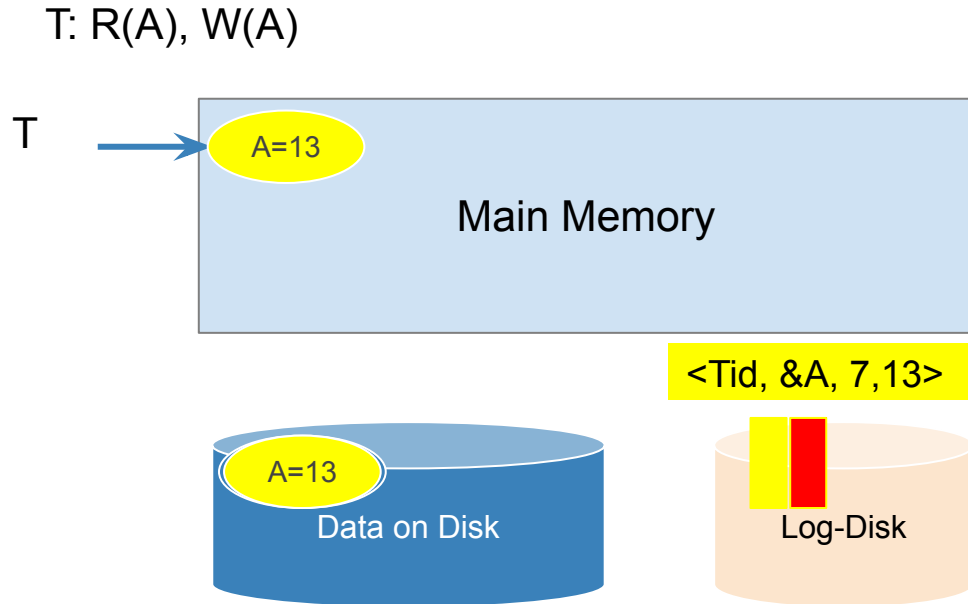


# Write-Ahead Logging (WAL) TXN Commit Protocol

# Write-ahead Logging (WAL) Commit Protocol



# Write-ahead Logging (WAL) Commit Protocol



Commit after we've written log to disk but before we've written data to disk... this is WAL!

**OK, Commit!**

If we crash now, is T durable?

**USE THE LOG!**

# Write-Ahead Logging (WAL)

## Algorithm: WAL

For each tuple update, write Update Record into LOG-RAM


Follow two **Flush** rules for LOG

- Rule1: **Flush** Update Record into LOG-Disk before corresponding data page goes to storage
- Rule2: Before TXN commits,
  - **Flush** all Update Records to LOG-Disk
  - **Flush** COMMIT Record to LOG-Disk


→ **Durability**

→ **Atomicity**

Transaction is committed *once COMMIT record is on stable storage*



Rule1: For each tuple update




Rule2: Before TXN commits



# Incorrect Commit Protocol #1

T: R(A), W(A)


A: 7 → 13



Let's try committing  
*before* we've written  
either data or LOG to  
disk...

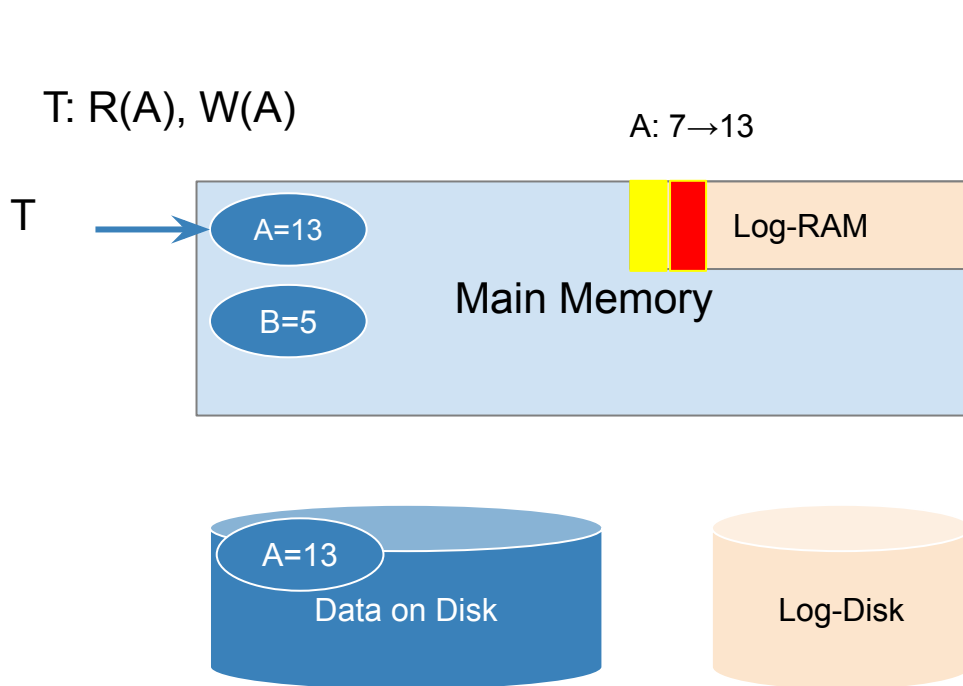
***OK, Commit!***

If we crash now, is T  
durable?



***Lost T's update!***

# Incorrect Commit Protocol #2



Let's try committing *after* we've written data but *before* we've written LOG to disk...

**OK, Commit!**

If we crash now, is T durable? Yes! Except...

**How do we know whether T was committed??**

# Example

## Monthly bank interest transaction

## Performance

Money

Account	....	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	....	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

WAL (@4:29 am day+1)

T-Monthly-423	<b>START TRANSACTION</b>		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352
T-Monthly-423	...	...	...
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	<b>COMMIT</b>		

### Cost to update all data

100M bank accounts → 100M seeks? (worst case)

(@10 msec/seek, that's 1 Million secs)



### Cost to Append to log

- + 1 seek to get 'end of log'
- + write 100M log entries sequentially (fast!!! < 10 sec)

[Lazily update data on disk later, when convenient.]

### Speedup for TXN Commit

1 Million secs vs 10 sec!!!

A close-up photograph of a person's hand holding a blue pen, poised to write on a white sheet of paper. The hand is wearing a grey, textured sweater. The background is blurred, showing a wooden desk and a laptop screen.

# Logging Summary

- If DB says TX commits, TX effect remains after database crash
- DB can undo actions and help us with atomicity
- This is only half the story...