



Intuition for SQL

Goals for

Today & Next
week

(SQL, SQL, SQL)

Phase I: Intuition for SQL (1st half of today)

Basic Relational model (aka tables)

Example SQL (exploring real datasets)

Phase II: Formal description

SQL concepts we'll study (similar to Python map-reduce)

Schemas, Query structure of SELECT-FROM-WHERE, JOINS, etc

A Motivating Example

A basic Course Management System (CMS):

Entities (e.g., Students, Courses, Professors)

Relationships (e.g., Who takes what, Who teaches what?)

Simple DB == Spreadsheets

The screenshot shows a Google Sheet interface with a menu bar (File, Edit, View, Insert, Format, Data, Tools, Add-ons, Help) and a status bar (All changes saved in Drive). The sheet contains two tables, each highlighted with a red border.

Students Table:

SID	Name	GPA
40001	Mickey	3.2
40002	Daffy	3.6
50003	Donald	3.3
50004	Minnie	3.9
10008	Pluto	4

Enrolled Table:

SID	CID	Grade
40001	1012	A
50004	1012	A-
40001	1017	B+

Courses Table:

CID	C-Name	Room
1012	CS145 - Toon systems	Nvidia
1017	CS161 - Animation art	Gates 300
1019	CS245 - Painting town red	Packard 45

Tables

Student(sid: *string*, name: *string*, gpa: *float*)

Courses(cid: *string*, c-name: *string*, room: *string*)

Enrolled(sid: *string*, cid: *string*, grade: *string*)
sid: Connects Enrolled to Students
cid: Connects Enrolled to Courses

Queries ["compute" over tables]

- Minnie's GPA?
- AVG student GPA?
- Mickey's classes?
- AVG student GPA in CS145?




Example SQL

Part I

Example 1

Sun Roof potential

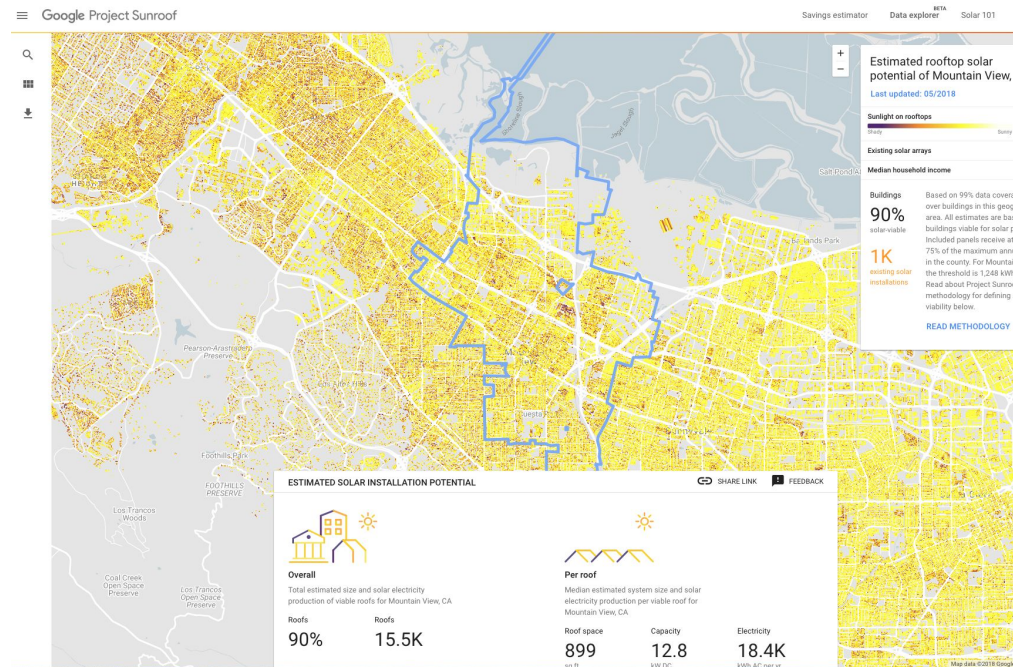
from Satellite
images



Example 1

SunRoof potential

SunRoof explorer



Example 1

Public Dataset: Solar potential by postal code

region_name	percent_covered	kw_total	carbon_offset_metric_tons
94043	97.79146031321109	215612.5	84929.00985071347
94041	99.05200433369447	56704.25	22189.34823862318

Public dataset

Public Dataset: USA.population by zip2010

zipcode	population
99776	124
38305	49808
37086	31513
41667	720
67001	1676

Example 1

SunRoof

On BigQuery
Public dataset

What is the solar potential of Mountain View, CA? [[Run query](#)]

🔗 Saved Query: MTV sunroof [edited] ?

```
1  #StandardSQL
2  SELECT
3    region_name,
4    percent_covered,
5    kw_total,
6    carbon_offset_metric_tons
7  FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code`
8  WHERE
9    region_name = '94040'
10   OR region_name = '94041'
11   OR region_name = '94043'
```

Ctrl + Enter: run

Standard SQL Dialect X

RUN QUERY ▼

Save Query

Save View

Format Query

Schedule Query

Show Options

Query complete (1.6s elapsed, 346 KB processed)

Results

Details

Download as CSV

Download as JSON

Save

Row	region_name	percent_covered	kw_total	carbon_offset_metric_tons
1	94043	97.79146031321109	215612.5	84929.00985071347
2	94041	99.05200433369447	56704.25	22189.34823862318
3	94040	98.9440337909187	139745.5	55039.74974407879

Example 2

SunRoof

Public dataset
On BigQuery

How many metric tons of carbon would we offset, if building in communities with 100% coverage all had solar roofs? [[Run query](#)]

🔗 Saved Query: CO2 offset in 100percent zips ⓘ

```
1 #StandardSQL
2 SELECT
3   ROUND(SUM(s.carbon_offset_metric_tons),2) total_carbon_offset_possible_metric_tons
4 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code` s
5 JOIN `bigquery-public-data.census_bureau_usa.population_by_zip_2010` c
6 ON s.region_name = c.zipcode
7 WHERE
8   percent_covered = 100.0
9   AND c.population > 0
10
11
12
13
```

Standard SQL Dialect ✕

Ctrl + Enter: run q

RUN QUERY ▾

Save Query

Save View

Format Query

Schedule Query

Show Options

Query com

Results

Details

Download as CSV

Download as JSON

Save as

Row	total_carbon_offset_possible_metric_tons
-----	--

1	3689508.33
---	------------

Example 2

SunRoof

Public dataset
On BigQuery

How many metric tons of carbon would we offset, per zipcode?

⇒ Saved Query: CO2 offset in 100percent zips [edited] ?

```
1 #StandardSQL
2 SELECT
3   zipcode, ROUND(SUM(s.carbon_offset_metric_tons),2) total_carbon_offset_possible_metric_tons
4 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code` s
5 JOIN `bigquery-public-data.census_bureau_usa.population_by_zip_2010` c
6 ON s.region_name = c.zipcode
7 WHERE
8   percent_covered = 100.0
9   AND c.population > 0
10 GROUP BY c.zipcode
11
12
13
```

Standard SQL Dialect ✕

Ctrl + Enter: run c

RUN QUERY ▼

Save Query

Save View

Format Query

Schedule Query

Show Options

Query complete (2.5s elapsed, 23.5 MB processed)

Results	Details	Download as CSV	Download as JSON	Save a
Row	zipcode	total_carbon_offset_possible_metric_tons		
1	35119	3417.26		
2	10165	162.1		
3	21810	6650.09		
4	74078	61515.66		
5	47876	5544.3		
6	10170	831.22		

Example 2

SunRoof

How many metric tons of carbon would we offset, per zipcode sorted?

Saved Query: CO2 offset in 100percent zips [edited]

Query Editor

```
1 #StandardSQL
2 SELECT
3   zipcode, ROUND(SUM(s.carbon_offset_metric_tons),2) total_carbon_offset_possible_metric_tons
4 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code` s
5 JOIN `bigquery-public-data.census_bureau_usa.population_by_zip_2010` c
6 ON s.region_name = c.zipcode
7 WHERE
8   percent_covered = 100.0
9   AND c.population > 0
10 GROUP BY c.zipcode
11 ORDER BY total_carbon_offset_possible_metric_tons
12 DESC
13
14
```

Standard SQL Dialect

Ctrl + Enter: run query, Tab or Ctrl + Space

RUN QUERY

Save Query

Save View

Format Query

Schedule Query

Show Options

Query complete (2.8s elapsed, 23.5 MB processed)

Results

Details

Download as CSV

Download as JSON

Save as Table

Save to

Row	zipcode	total_carbon_offset_possible_metric_tons
1	18503	715700.55
2	44243	271861.55
3	38677	266787.12
4	96860	225850.35
5	47809	141087.91

Query with SQL, universally over 'all' DBs

Reminder

Special
Databases






SQL

```
= QUERY(T,  
  "SELECT c1, c2  
   FROM T  
  WHERE condition;")
```

```
SELECT c1, c2  
FROM T  
WHERE condition;
```

```
results =  
spark.SQL(  
  "SELECT c1, c2  
   FROM T  
  WHERE condition;")
```

DB



'Spreadsheets'

GCP BigQuery, AWS Redshift,
MySQL, PostgreSQL, Oracle

Spark, Hadoop

Data

Data

100s of Scaling algorithms/systems? [Weeks 3..]

- Data layout? [Row vs columns...]
- Data structs? [Indexing...]

Preview

SQL queries

QUERYING DATA FROM A TABLE

SELECT c1, c2 FROM t;
Query data in columns c1, c2 from a table

SELECT * FROM t;
Query all rows and columns from a table

SELECT c1, c2 FROM t
WHERE condition;
Query data and filter rows with a condition

SELECT DISTINCT c1 FROM t
WHERE condition;
Query distinct rows from a table

SELECT c1, c2 FROM t
ORDER BY c1 ASC [DESC];
Sort the result set in ascending or descending order

SELECT c1, c2 FROM t
ORDER BY c1
LIMIT n OFFSET offset;
Skip *offset* of rows and return the next *n* rows

SELECT c1, aggregate(c2)
FROM t
GROUP BY c1;
Group rows using an aggregate function

SELECT c1, aggregate(c2)
FROM t
GROUP BY c1
HAVING condition;
Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

SELECT c1, c2
FROM t1
INNER JOIN t2 ON condition;
Inner join t1 and t2

SELECT c1, c2
FROM t1
LEFT JOIN t2 ON condition;
Left join t1 and t2

SELECT c1, c2
FROM t1
RIGHT JOIN t2 ON condition;
Right join t1 and t2

SELECT c1, c2
FROM t1
FULL OUTER JOIN t2 ON condition;
Perform full outer join

SELECT c1, c2
FROM t1
CROSS JOIN t2;
Produce a Cartesian product of rows in tables

SELECT c1, c2
FROM t1, t2;
Another way to perform cross join

SELECT c1, c2
FROM t1 A
INNER JOIN t2 B ON condition;
Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT c1, c2 FROM t1
UNION [ALL]
SELECT c1, c2 FROM t2;
Combine rows from two queries

SELECT c1, c2 FROM t1
INTERSECT
SELECT c1, c2 FROM t2;
Return the intersection of two queries

SELECT c1, c2 FROM t1
MINUS
SELECT c1, c2 FROM t2;
Subtract a result set from another result set

SELECT c1, c2 FROM t1
WHERE c1 [NOT] LIKE pattern;
Query rows using pattern matching %, _

SELECT c1, c2 FROM t
WHERE c1 [NOT] IN value_list;
Query rows in a list

SELECT c1, c2 FROM t
WHERE c1 BETWEEN low AND high;
Query rows between two values

SELECT c1, c2 FROM t
WHERE c1 IS [NOT] NULL;
Check if values in a table is NULL or not